



Views (Virtual Tables) in SQL

Database Design

Department of Computer Engineering

Sharif University of Technology

Maryam Ramezani maryam.ramezani@sharif.edu

Introduction

Concept of view



- ❑ We may frequently issue queries that retrieve the employee name and the project names that the employee works on. Rather than having to specify the join of the three tables EMPLOYEE, WORKS_ON, and PROJECT every time we issue this query, we can define a view that is specified as the result of these joins. Then we can issue queries on the view, which are specified as single table retrievals rather than as retrievals involving two joins on three tables. We call the EMPLOYEE, WORKS_ON, and PROJECT tables the defining tables of the view.

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------

DEPT_LOCATIONS

<u>Dnumber</u>	<u>Dlocation</u>
----------------	------------------

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

DEPENDENT

<u>Essn</u>	<u>Dependent_name</u>	Sex	Bdate	Relationship
-------------	-----------------------	-----	-------	--------------



- ❑ A **view** in SQL terminology is a single table that is derived from other tables.
- ❑ These other tables can be *base tables* or previously defined views.
- ❑ A view does not necessarily exist in physical form; it is considered to be a **virtual table**, in contrast to **base tables**, whose tuples are always physically stored in the database.
 - This limits the possible update operations that can be applied to views, but it does not provide any limitations on querying a view.
- ❑ We can think of a view as a way of specifying a table that we need to reference frequently, even though it may not exist physically.
- ❑ The view is given a (virtual) table name (or view name), **a list of attribute names, and a query to specify the contents of the view.**
- ❑ If none of the view attributes results from applying functions or arithmetic operations, we do not have to specify new attribute names for the view, since they would be the same as the names of the attributes of the defining tables in the default case.

Specification of Views in SQL



CREATE VIEW
AS SELECT
FROM
WHERE

WORKS_ON1
Fname, Lname, Pname, Hours
EMPLOYEE, PROJECT, WORKS_ON
Ssn = Essn **AND** Pno = Pnumber;

WORKS_ON1

Fname	Lname	Pname	Hours
-------	-------	-------	-------

SELECT Fname, Lname
FROM WORKS_ON1
WHERE Pname = 'ProductX';

DROP VIEW WORKS_ON1;

CREATE VIEW
AS SELECT
FROM
WHERE
GROUP BY

DEPT_INFO(Dept_name, No_of_emps, Total_sal)
Dname, **COUNT** (*), **SUM** (Salary)
DEPARTMENT, EMPLOYEE
Dnumber = Dno
Dname;

DEPT_INFO

Dept_name	No_of_emps	Total_sal
-----------	------------	-----------



- ❑ SQL query on a view is in the same way we specify queries on base tables:

- Example: retrieve the last name and first name of all employees who work on the 'ProductX' project
- QV1:

```
SELECT Fname, Lname  
FROM WORKS_ON1  
WHERE Pname = 'ProductX';
```

WORKS_ON1

Fname	Lname	Pname	Hours
-------	-------	-------	-------



❑ Advantages:

- Simplify the specification of certain queries.
- Used as a security and authorization mechanism (End of this slide!)

❑ A view is supposed to be always up-to-date; if we modify the tuples in the base tables on which the view is defined, the view must automatically reflect these changes. Hence, **the view does not have to be realized or materialized at the time of view definition but rather at the time when we specify a query on the view.**

❑ It is the responsibility of the DBMS and not the user to make sure that the view is kept up-to-date.



- ❑ If we do not need a view anymore, we can use the **DROP VIEW** command to dispose of it.
 - Example: DROP VIEW WORKS_ON1

View Implementation in DBMs



- How a DBMS can efficiently implement a view for efficient querying?
 - It is complex!
 - Two main approaches have been suggested:
 - Query modification
 - View materialization



- ❑ Involves modifying or transforming the view query (submitted by the user) into a query on the underlying base tables.

CREATE VIEW	WORKS_ON1	QV1:	SELECT	Fname, Lname
AS SELECT	Fname, Lname, Pname, Hours		FROM	WORKS_ON1
FROM	EMPLOYEE, PROJECT, WORKS_ON		WHERE	Pname = 'ProductX';
WHERE	Ssn = Essn AND Pno = Pnumber;			

SELECT	Fname, Lname
FROM	EMPLOYEE, PROJECT, WORKS_ON
WHERE	Ssn = Essn AND Pno = Pnumber
	AND Pname = 'ProductX';

Query modification (View Computation)



EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

SELECT
FROM
WHERE

Fname, Lname
EMPLOYEE, PROJECT, WORKS_ON
Ssn = Essn **AND** Pno = Pnumber
AND Pname = 'ProductX';

- **Disadvantage:**
Inefficient for views defined via complex queries that are time-consuming to execute, especially if multiple view queries are going to be applied to the same view within a short period of time



- How a DBMS can efficiently implement a view for efficient querying?
 - It is complex!
 - Two main approaches have been suggested:
 - Query modification
 - View materialization



- ❑ Involves physically creating a temporary or permanent view table when the view is first queried or created and keeping that table on the assumption that other queries on the view will follow.
- ❑ Efficient strategy for automatically updating the view table when the base tables are updated? **incremental update** have been developed for this purpose, where the DBMS can determine what new tuples must be inserted, deleted, or modified in a materialized view table when a database update is applied to one of the defining base tables.



- ❑ The view is generally kept as a materialized (physically stored) table as long as it is being queried.
- ❑ If the view is not queried for a certain period of time, the system may then automatically remove the physical table and recompute it from scratch when future queries reference the view.
- ❑ Application: For Data mining goals!

Materialized View Example



```
CREATE MATERIALIZED VIEW user_purchase_summary AS SELECT
    u.id as user_id,
    COUNT(*) as total_purchases,
    SUM(CASE when p.status = 'cancelled' THEN 1 ELSE 0 END) as cancelled_purchases
FROM users u
JOIN purchases p ON p.user_id = u.id;
```

When executed, this statement instructs the database to:

- Execute the SELECT query within the materialized view definition.
- Cache the results in a new “virtual” table named user_purchase_summary
- Save the original query so it knows how to update the materialized view in the future.



- ❑ Different strategies as to when a materialized view is updated are possible:
 - **Immediate update** strategy updates a view as soon as the base tables are changed
 - **Lazy update** strategy updates the view when needed by a view query
 - **Periodic update** strategy updates the view periodically (in the latter strategy, a view query may get a result that is not up-to-date).

How do materialized views work in specific databases?



Database	Materialized Views?	View Maintenance	Notes
PostgreSQL	Yes, in v9.3+	Manual	Materialized views are populated at time of creation and must be manually refreshed via REFRESH MATERIALIZED VIEW statements that recompute the entire view.
MySQL	No	N/A	
Microsoft SQL Server	Yes	Automatic	SQL Server calls them "Indexed Views" because the materialization step is a matter of creating an index on a regular view. SQL Server limits indexed views to basic SQL queries .
Oracle	Yes	Multiple Options	Materialized views in Oracle databases can be set to manually refresh, refresh on a schedule, or, if the SQL query meets these requirements , automatically refreshed.

User Commands on View



- ❑ A user can always issue a retrieval query against any view.
- ❑ Issuing an INSERT, DELETE, or UPDATE command on a view table is in **many cases not possible**.
- ❑ An update on a view defined on a single table without any aggregate functions can be mapped to an update on the underlying base table under certain conditions.
- ❑ For a view involving joins, an update operation may be mapped to update operations on the underlying base relations in multiple ways. Hence, it is often not possible for the DBMS to determine which of the updates is intended.



- ❑ **Updatable**: Can be updated but has problems in some cases.
- ❑ **Non Updatable**: Can not convert external to conceptual.

- ❑ View can be defined on:
 - **One base table**:
 - Key Preserving: view has primary key of base table: OK with some problems.
 - Non Key Preserving: NO
 - View has aggregated columns: NO

 - **More than one base table**: In practice non updatable but updatable in theory.

Example



EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------

WORKS_ON

<u>Essn</u>	<u>Pno</u>	Hours
-------------	------------	-------

**CREATE VIEW
AS SELECT
FROM
WHERE**

WORKS_ON1

Fname, Lname, Pname, Hours
EMPLOYEE, PROJECT, WORKS_ON
Ssn = Essn **AND** Pno = Pnumber;

WORKS_ON1

Fname	Lname	Pname	Hours
-------	-------	-------	-------



- ❑ Challenge 1: Consider the WORKS_ON1 view, and suppose that we issue the command to update the PNAME attribute of 'John Smith' from 'ProductX' to 'ProductY'.

```
UV1:      UPDATE WORKS_ON1
          SET      Pname = 'ProductY'
          WHERE    Lname = 'Smith' AND Fname = 'John'
               AND Pname = 'ProductX';
```

Updating a view defined on multiple tables



- ❑ Two possible updates on the base relations corresponding to the view update operation in UV1:

(a): **UPDATE WORKS_ON**
SET Pno = (**SELECT** Pnumber
FROM PROJECT
WHERE Pname = 'ProductY')
WHERE Essn IN (**SELECT** Ssn
FROM EMPLOYEE
WHERE Lname = 'Smith' **AND** Fname = 'John')
AND
Pno = (**SELECT** Pnumber
FROM PROJECT
WHERE Pname = 'ProductX');

relates 'John Smith' to the
'ProductY' PROJECT tuple
instead of the
'ProductX' PROJECT tuple
and is the most likely desired
update.

Updating a view defined on multiple tables



- ❑ Two possible updates on the base relations corresponding to the view update operation in UV1:

(b): **UPDATE** PROJECT **SET** Pname = 'ProductY'
 WHERE Pname = 'ProductX';

(b) Would also give the desired update effect on the view, but it accomplishes this by changing the name of the 'ProductX' tuple in the PROJECT relation to 'ProductY'. It is quite unlikely that the user who specified the view update UV1 wants the update to be interpreted as in (b), since it also has the side effect of changing all the view tuples with Pname = 'ProductX'.

Updating a view defined on multiple tables



- ❑ Challenge 2: Some view updates may not make much sense; for example, modifying the Total_sal attribute of the DEPT_INFO view does not make sense because Total_sal is defined to be the sum of the individual employee salaries. This incorrect request is shown as UV2:

```
V2:      CREATE VIEW      DEPT_INFO(Dept_name, No_of_emps, Total_sal)
        AS SELECT        Dname, COUNT (*), SUM (Salary)
        FROM              DEPARTMENT, EMPLOYEE
        WHERE              Dnumber = Dno
        GROUP BY           Dname;
```

```
UV2:      UPDATE      DEPT_INFO
        SET            Total_sal = 100000
        WHERE           Dname = 'Research';
```



- ❑ In SQL, the clause **WITH CHECK OPTION** should be added at the end of the view definition if a view is to be updated by INSERT, DELETE, or UPDATE statements.

```
CREATE VIEW V2 [(SN, SJ, SL)]  
AS SELECT STID, STJ, STL  
FROM STT  
WHERE STJ != 'phys'  
[WITH CHECK OPTION]
```

- ❑ This allows the system to reject operations that violate the SQL rules for view updates. The full set of SQL rules for when a view may be modified by the user are more complex than the rules stated earlier.

View Example



Diagram illustrating a view definition. A blue arrow points from the **STID** column of the **STT** table to the **SN** column of the **V2** view.

STT	STID	STNAME	STLEV	STMJR	STD
	777	st7	bs	phys	d11
	888	st8	ms	math	d12
	444	st4	bs	comp	d14
	⋮	⋮	⋮	⋮	⋮

V2	SN	SL	SJ
	888	ms	math
	444	bs	comp
	⋮	⋮	

```
CREATE VIEW V2 [(SN, SJ, SL)]
AS SELECT STID, STMJR, STLEV
FROM STT
WHERE STMJR != 'phys'
[WITH CHECK OPTION]
```



```
DELETE FROM V2  
WHERE SN='444'
```



```
DELETE FROM STT  
WHERE STID='444' AND STMJR != 'phys'
```

- ❑ Any single row operation is ok! This is a single row because we have **STID** the primary key in the condition of query!
- ❑ The row is deleted, and any other user will not have this row in his view!



```
DELETE FROM V2  
WHERE SL='ms'
```



```
DELETE FROM STT  
WHERE STLEV='ms' AND STMJR != 'phys'
```

- ❑ This is not a single row, because we don't have primary key in query condition. Many rows in STT may be deleted.
- ❑ The request is rejected by the DBMS!



CREATE VIEW V3

**AS SELECT DISTINCT STNAME, STMJR
FROM STT**



DELETE FROM V3

WHERE STNAME='a'



Rejected by the DBMS!



```
UPDATE V2  
SET SJ='IT'  
WHERE SN='444'
```



```
UPDATE STT  
SET STMJR='IT'  
WHERE STID='444' ' AND STMJR != 'phys'
```




UPDATE V2

SET SJ='phys'

WHERE SN='888'



UPDATE STT

SET STMJR='phys'

WHERE STID='888' ' AND STMJR != 'phys'



Rejected because of **WITH CHECK OPTION**

If there was not **WITH CHECK OPTION**, what happened?

SELECT V2.* FROM V2

No row with 888 key in view! :D

Summarized updating a view by user



- ❑ Generally, a view update is feasible **when only one possible update** on the base relations can accomplish the desired update operation on the view.
- ❑ Whenever an update on the view can be mapped **to more than one update on the underlying base relations**, it is usually not permitted.

Note: Some researchers have suggested that the DBMS have a certain procedure for choosing one of the possible updates as the most likely one. Some researchers have developed methods for choosing the most likely update, whereas other researchers prefer to have the user choose the desired update mapping during view definition. **But these options are generally not available in most commercial DBMSs**



INSERT INTO V2

VALUES ('555', 'chem', 'bs')

INSERT INTO STT

VALUES ('555', ?, 'chem', 'bs', ?)

- ❑ The request is rejected, if the hidden columns has NOT NULL constraint.
- ❑ If STID is a primary key and '555' was inserted in STT table before but is hidden in view, then the request is rejected.
- ❑ Inserting value **VALUES** ('555' , 'phys' , 'bs') is rejected because of **[WITH CHECK OPTION]**.

View with Aggregated Columns



- ❑ It is non updatable in practice and theory

V4	PN	SQ
	P1	100
	P2	210
	P3	80

```
CREATE VIEW V4 (PN, SQ)
AS SELECT P#, SUM(QTY)
FROM SP
GROUP BY P#
```

SP	S#	P#	QTY
	S1	P1	100
	S1	P2	140
	S2	P3	80
	S2	P2	70

```
DELETE FROM V4
WHERE PN='p2'
```





A user **can not** insert data into view with a single defining table if:

- ❑ Having **[WITH CHECK OPTION]** and conflict with the constraints.
- ❑ Violation of the **UNIQUENESS** constraint of primary key.
- ❑ Violation of the **NOT NULL** constraint



- ❑ A view with a single defining table is updatable if:
 - (1) the view attributes contain the primary key of the base relation
 - (2) All the other attributes of base table do not have the NOT NULL constraint or has the default values
 - (3) View query does not have Distinct in the select.
 - (4) View query does not have group by or having.

- ❑ Views defined on multiple tables using joins are generally not updatable.

- ❑ Views defined using grouping and aggregate functions are not updatable.

Theory:

Update View of Multiple Table



- ❑ As practically view of multiple base tables is not updatable, but when is it theoretically acceptable?
 - Cartesian Product ✗
 - Join on two primary keys PK-PK
 - Join on primary and foreign keys PK-FK
 - Join on two foreign keys which are referenced to primary key of another table FK-FK
 - Join on two non-key columns of two tables without any relationship NK-NK
 - Union
 - Intersect
 - Except

Multiple Table Example



V5

STID	STNAME	STLEV	STMJR	STDEID
777	st7	bs	phys	d11
888	st8	ms	math	d12
444	st4	bs	comp	d14
:	:	:	:	:

```
CREATE VIEW V5
AS SELECT *
FROM ST1 NATURAL JOIN ST2
```

ST1

STID	STNAME	STLEV
777	st7	bs
888	st8	ms
444	st4	bs
:	:	:

ST2

STID	STMJR	STDEID
777	phys	d11
888	math	d12
444	comp	d14
:	:	:



INSERT INTO V5

VALUES ('999', 'St9', 'chem', 'bs', 'D15')

INSERT INTO ST1

VALUES ('999', 'St9', 'bs')

INSERT INTO ST2

VALUES ('999', 'chem', 'D15')

Multiple Table PK–FK



STT

STID	STNAME	STLEV	STMJR	STD
777	st7	bs	phys	d11
888	st8	ms	math	d12
444	st4	bs	comp	d14
⋮	⋮	⋮	⋮	⋮

STCOT

STID	COURSE	GRADE
777	40280	19
888	40567	20
444	40232	15
⋮	⋮	⋮



CREATE VIEW V6

AS SELECT STT.STID, STT.NAME, STCOT.COURSE,STCOT.GRADE
FROM STT **JOIN** STCOT

INSERT INTO V6

VALUES ('9212345', 'Amir', '40638', 15)



INSERT INTO STT

VALUES ('9212345', 'Amir', ?, ?, ?)

INSERT INTO STCOT

VALUES ('9212345', '40638', 15)

***Insert one row in STCOT (table having foreign key!)**

***If the student is not in table STT then it will be inserted.**



CREATE VIEW V6

AS SELECT STT.STID, STT.NAME, STCOT.COURSE,STCOT.GRADE
FROM STT **JOIN** STCOT

Delete from V6

where STID='9212345' and NAME= 'Amir' and STCOT= '40638' and GRADE=15)



Delete from STCOT

What about STT? Affect on other tables and view?



Create view v7 as select * from T1 union select * from T2

- ☐ Insert row: into both or one of T1 or T2.
- ☐ Delete row: delete from both T1 and T2.
- ☐ Update row: update both in T1 and T2.



Create view v8 as select * from T1 intersect select * from T2

- ☐ Insert row: insert into both T1 and T2.
- ☐ Delete row: delete from T1 or T2.
- ☐ Update row: update both in T1 and T2.



Create view v9 as select * from T1 except select * from T2

- ☐ Insert row: insert into T1 and check not be in T2.
- ☐ Delete row: delete from T1
- ☐ Update row: update in T1 and check not be in T2.

In-line View



- ❑ It is also possible to define a view table in the FROM clause of an SQL query. This is known as an in-line view. In this case, the view is defined within the query itself.
- ❑ **Inline views** refer to a SELECT statement located in the FROM clause of secondary SELECT statement. Inline views can help make complex queries simpler by removing compound calculations or eliminating join operations while condensing several separate queries into a single simplified query.



- ❑ PostgreSQL semantics may refer to inline views as Subselect or as Subquery
- ❑ ERROR: subquery in FROM must have an alias. This is because in PostgreSQL the use of aliases is mandatory. The following example uses B as an alias.

```
SELECT A.LAST_NAME, A.SALARY, A.DEPARTMENT_ID, B.SAL_AVG
FROM EMPLOYEES A,
  (SELECT DEPARTMENT_ID, ROUND(AVG(SALARY)) AS SAL_AVG
   FROM EMPLOYEES GROUP BY DEPARTMENT_ID) B
WHERE A.DEPARTMENT_ID = B.DEPARTMENT_ID;
```

Authorization Mechanisms with View



- ❑ Creating an appropriate view and granting certain users access to the view and not the base tables, they would be restricted to retrieving only the data specified in the view.

- Access to specific rows

```
CREATE VIEW    DEPT5EMP    AS
SELECT        *
FROM          EMPLOYEE
WHERE         Dno = 5;
```

- Access to specific columns

```
CREATE VIEW    BASIC_EMP_DATA    AS
SELECT        Fname, Lname, Address
FROM          EMPLOYEE;
```



❑ Advantages:

- Simplify the specification of certain queries.
- Used as a security and authorization mechanism (for columns & rows)
 - Data sharing
 - Hidden Data
- Data Independence (Next Topic!)
- Good for data mining and reporting



❑ Disadvantages:

- Not good when system is single user!
- Not good when user needs to run updatable commands.
- Overhead for converting query to base table queries